

Report for Salary Dataset Simple Linear Regression

1. Dataset Introduction

Salary Dataset in CSV for Simple linear regression.

Columns:

#

YearsExperience

Salary

Data source: <https://www.kaggle.com/datasets/abhishek14398/salary-dataset-simple-linear-regression/data>

2. Descriptive Statistics

shape: (9, 4)

statistic		YearsExperience	Salary
---	---	---	---
str	f64	f64	f64
count	30.0	30.0	30.0
null_count	0.0	0.0	0.0
mean	14.5	5.413333	76004.0
std	8.803408	2.837888	27414.429785
min	0.0	1.2	37732.0
25%	7.0	3.3	56643.0
50%	15.0	5.0	66030.0
75%	22.0	8.0	101303.0
max	29.0	10.6	122392.0

3. Profiler benchmark for Polars vs Pandas

I use Profiler from pyinstrument to compare Polars & Pandas.

For my dataset, Polars took 0.001s while Pandas took 0.012s.

Polars generally outperforms Pandas in terms of speed, especially for large datasets. This is because Polars is designed for parallelized execution and is optimized for in-memory performance.

Pandas may still be more familiar or convenient for certain smaller tasks or when using legacy systems that require it.

(30, 3)
shape: (9, 4)

statistic		YearsExperience	Salary
str	f64	f64	f64
count	30.0	30.0	30.0
null_count	0.0	0.0	0.0
mean	14.5	5.413333	76004.0
std	8.803408	2.837888	27414.429785
min	0.0	1.2	37732.0
25%	7.0	3.3	56643.0
50%	15.0	5.0	66030.0
75%	22.0	8.0	101303.0
max	29.0	10.6	122392.0

Recorded: 15:25:00 Samples: 0
Duration: 0.001 CPU time: 0.002
v4.7.3

Profile at /Users/xianjinghuang/Desktop/Xianjing_Huang_Mini_Proj_3/test_main.py:46

No samples were recorded.

(30, 3)

	Unnamed: 0	YearsExperience	Salary
count	30.000000	30.000000	30.000000
mean	14.500000	5.413333	76004.000000
std	8.803408	2.837888	27414.429785
min	0.000000	1.200000	37732.000000
25%	7.250000	3.300000	56721.750000
50%	14.500000	4.800000	65238.000000
75%	21.750000	7.800000	100545.750000
max	29.000000	10.600000	122392.000000

Recorded: 15:25:00 Samples: 0
Duration: 0.012 CPU time: 0.008
v4.7.3

4. Data Visualization

